

# OpenADR 3.0 Demand Flexibility for Heat Pump Water Heaters

IEA EBC Annex 96 – Activity A3

**Anand Prakash**

Carnegie Mellon University

[anandkrp@andrew.cmu.edu](mailto:anandkrp@andrew.cmu.edu)

# Outline

- 1. Introduction** — Motivation and goals
- 2. Background** — OpenADR 3.0, demand flexibility, HPWHs
- 3. Use Case** — Price-responsive HPWH control through OpenADR
- 4. Software Architecture** — Components and data flow
- 5. Implementation** — LP/heuristic schedulers and CTA-2045 integration
- 6. How to Use** — Setup, run, and extend

# 1. Introduction

## Why Demand Flexibility?

- Grid decarbonization requires flexible loads that can **shift consumption**
- Water heating accounts for **~18% of residential energy use** in the US
- Heat pump water heaters (HPWHs) with storage tanks are ideal candidates:
  - Thermal storage enables **load shifting** without impacting comfort
  - Can pre-heat during cheap/clean hours, coast during expensive/dirty hours
- **Challenge:** How do we communicate price signals to devices at scale?

# Project Goals

Develop an **open-source software toolkit** so that researchers and practitioners can:

1. Set up an OpenADR 3.0 communication infrastructure (VTN + VEN)
2. Fetch real electricity pricing data and publish it as OpenADR events
3. Run a control algorithm that converts price signals into HPWH schedules
4. Generate CTA-2045 demand response commands for water heaters
5. Test the full pipeline end-to-end on their own machines

All code, documentation, and quickstart notebooks are publicly available.

## 2. Background

# OpenADR 3.0

**Open Automated Demand Response** — an open standard for communicating DR signals.

Concept	Description
<b>VTN</b> (Virtual Top Node)	Server — publishes programs, events, price signals
<b>VEN</b> (Virtual End Node)	Client — receives signals, controls devices
<b>Program</b>	Defines a demand response program (e.g., dynamic pricing)
<b>Event</b>	Time-based signal with payload (e.g., hourly prices)
<b>Report</b>	Telemetry data sent from VEN back to VTN

REST API (JSON over HTTP) with OAuth 2.0 authentication

# Heat Pump Water Heaters as Flexible Loads

## How HPWHs provide flexibility:

- Tank stores thermal energy (50–80 gal)
- Heat pump COP of 3–4x vs resistance
- Can **pre-heat** during low-price hours, **coast** during high-price hours
- No comfort impact if managed well

Parameter	Typical Value
Tank capacity	50–80 gallons
HP output	4–5 kW thermal
COP	2.5–4.5
Thermal storage	8–15 kWh

# CTA-2045 and Device Communication

**CTA-2045** is a modular communications interface for energy devices, providing standardized demand response commands for water heaters:


Signal	Code	Water Heater Action
<b>Shed</b>	-1	Lower setpoint, disable HP — coast on stored energy
<b>Normal</b>	0	Default operation
<b>Load Up</b>	1	Raise setpoint, pre-heat the tank
<b>Adv. Load Up</b>	2	Max setpoint, tight deadband

This project covers the full pipeline: **OpenADR → Control Algorithm → CTA-2045**

### **3. Use Case**


# Price-Responsive HPWH Control Through OpenADR

**Scenario:** A utility publishes dynamic electricity prices. A water heater controller receives these prices and optimizes its operation schedule.



OpenADR 3.0 Price Signal → LP Scheduler → CTA-2045 Commands → Device Action


# Example: 24-Hour Price Signal and Response



**Result:** Same hot water delivered at lower cost by shifting to cheap hours.

## **4. Software Architecture**

# System Architecture



# Repository Structure

```
annex96-a3-hotwater/
├── README.md
├── requirements.txt
├── instructions.ipynb
├── instructions-openleadr.ipynb
├── quickstart.ipynb
├── quickstart-openleadr.ipynb
└── controls/
    ├── hpwh_load_shift_lp.py
    ├── hpwh_load_shift_heuristic.py
    └── cta2045.py
└── sample_data/
└── presentation/
```

# Project overview  
# Python dependencies  
# Setup: Python VTN  
# Setup: openleadr-rs VTN  
# Demo: Python VTN  
# Demo: openleadr-rs VTN  
# Control algorithms  
# LP scheduler (globally optimal)  
# Heuristic scheduler (greedy)  
# CTA-2045 schedule generation  
# Example JSON payloads  
# This presentation

## Two VTN Options

	Python VTN Ref. Impl.	openleadr-rs (Rust)
<b>Language</b>	Python (Flask)	Rust (Axum)
<b>Database</b>	In-memory	PostgreSQL (Docker)
<b>Auth</b>	bl_client/1001	any-business
<b>Base URL</b>	localhost:8080/openadr3/3.0.1	localhost:3000
<b>Access</b>	Contact for access	Open source
<b>Best for</b>	Quick local testing	Production-like setup


## 5. Implementation

# HPWH Load Shift Scheduler

Two interchangeable implementations in `controls/`:

	<b>LP Scheduler ( <code>hphw_load_shift_lp</code> )</b>	<b>Heuristic ( <code>hphw_load_shift_heuristic</code> )</b>
<b>Method</b>	Linear program (HiGHS via scipy)	Bottom-up greedy
<b>Solution</b>	Globally optimal	Near-optimal
<b>Speed</b>	Fast (milliseconds)	$O(N^2)$ worst case
<b>Dep.</b>	scipy	numpy only

**LP formulation:**  $\min \sum e[h] \cdot \text{price}[h]/\text{COP}[h]$  subject to storage bounds and HP capacity bounds.



# CTA-2045 Schedule Generation



Full pipeline: OpenADR price signal → LP-optimal schedule → CTA-2045 commands

Two approaches: from LP/heuristic scheduler output (uses HP output levels), or directly from prices (uses percentile thresholds).

# Quickstart Demo Pipeline

```
# Step 1: Authenticate with VTN
token = get_token("bl_client", "1001")

# Step 2: Fetch prices from Olivine API
olivine_data = requests.get(OLIVINE_PRICING_URL).json()

# Step 3: Create pricing program on VTN
requests.post(f"{VTN}/programs", json=program_data, headers=auth)

# Step 4: Publish price event
requests.post(f"{VTN}/events", json=event_data, headers=auth)

# Step 5: Read events as VEN
events = requests.get(f"{VTN}/events", headers=ven_auth).json()

# Step 6: Run LP Scheduler → Step 7: Generate CTA-2045 schedule
schedule, converged = hpwh_load_shift(params)
cta_schedule = hpwh_load_shift_to_cta2045(schedule, params)
```

# Example Output

```
LP status: Optimization terminated successfully (HiGHS Status 7: Optimal)
```

```
Hourly schedule (kWh):
```

```
Hour 0: OFF 0.00 kWh @ $0.12052/kWh → Shed
Hour 1: OFF 0.00 kWh @ $0.12227/kWh → Shed
...
Hour 18: ON 1.50 kWh @ $0.11120/kWh → Normal
Hour 19: ON 1.50 kWh @ $0.10689/kWh → Normal
Hour 20: ON 1.50 kWh @ $0.10519/kWh → Normal
Hour 21: ON 1.50 kWh @ $0.10300/kWh → Normal
Hour 22: ON 1.50 kWh @ $0.10620/kWh → Normal
Hour 23: ON 1.50 kWh @ $0.10930/kWh → Normal
```

```
Total electricity cost: $0.28741
```

LP finds the globally optimal allocation — charges exactly what is needed at the cheapest hours.

# **Quickstart: Step by Step**

Each step – inputs, what it does, outputs, and what to change for your implementation

# Step 1: Setup & Verify VTN Connection

## 1 Setup & Verify VTN Connection


Exchange credentials for an OAuth token; confirm the VTN is reachable

INPUTS	WHAT IT DOES	OUTPUTS
<ul style="list-style-type: none"><li>• <b>VTN_BASE_URL</b> localhost:8080/openadr3/3.0.1</li><li>• <b>BL_CLIENT_ID / SECRET</b> Business Logic credentials</li><li>• <b>VEN_CLIENT_ID / SECRET</b> VEN read-only credentials</li></ul>	<p><b>POST /auth/token</b> grant_type: client_credentials → returns JWT Bearer token</p> <p><b>GET /programs</b> Confirm VTN returns HTTP 200 (raises exception if offline)</p>	<ul style="list-style-type: none"><li>• <b>access_token</b> Bearer JWT for API headers</li><li>• <b>bl_headers() helper</b> Authorization header for BL calls</li><li>• <b>ven_headers() helper</b> Authorization header for VEN calls</li><li>• <b>Confirmed VTN is live</b> exception raised if not running</li></ul>

**CUSTOMIZE**

- VTN implementation: change VTN\_BASE\_URL to localhost:3000 (no path prefix) for openleadr-rs
- Credentials: update CLIENT\_ID / SECRET to match your VTN's user fixture (see instructions.ipynb)
- Token caching: get\_token() is called per request; cache the token for higher-throughput use
- TLS: replace http:// with https:// and pass verify= cert path for non-local deployments

# Step 2: Fetch Dynamic Prices from Olivine



# Step 3: Create Pricing Program on VTN

**3 Create Pricing Program on VTN**

Register a program definition that groups related price events on the VTN

INPUTS	WHAT IT DOES	OUTPUTS
<ul style="list-style-type: none"><li>• OAuth token (Step 1) via bl_headers()</li><li>• <b>programName</b> "etou-dynamic-pricing"</li><li>• <b>payloadDescriptors</b> [{"payloadType": PRICE, "units": KWH}]</li></ul>	<p><b>POST /programs</b></p> <p>Register program definition on VTN VTN assigns a unique program_id</p> <p>-----</p> <p>program_id links all future events to this program (Step 4)</p> <p>VENs can filter events by program_id to subscribe to specific tariffs</p>	<ul style="list-style-type: none"><li>• <b>program_id</b> "0" (Python VTN) or UUID (openleadr-rs)</li><li>• <b>Stored program record</b> persists on VTN until deleted</li><li>• <b>Confirmed HTTP 201</b> creation verified via raise_for_status()</li></ul>

**CUSTOMIZE**

→ programName: any descriptive string matching your tariff (e.g. "eTOU-Dynamic", "real-time-pricing")  
→ Payload type: PRICE (absolute \$/kWh) or PRICE\_RELATIVE (delta from a baseline price)  
→ openleadr-rs extras: add programType, retailerName, country, principalSubdivision fields  
→ Reuse: query GET /programs first and skip creation if a matching program already exists

# Step 4: Publish Price Signal as an Event

**4** **Publish Price Signal as an Event**

Package the fetched prices as an OpenADR 3 event and post it to the VTN

INPUTS	WHAT IT DOES	OUTPUTS
<ul style="list-style-type: none"><li>• <b>prices + start_time (Step 2)</b> hourly prices and window start</li><li>• <b>program_id (Step 3)</b> links event to program</li><li>• <b>OAuth token (Step 1)</b> via bl_headers()</li></ul>	<p>Format each hour as an interval:</p> <pre>{id: hour, payloads: [{type: "PRICE", values: [price]}]}</pre> <p><b>POST /events</b></p> <p>Event stored on VTN server Any auth'd VEN can now read it via GET /events (Step 5)</p>	<ul style="list-style-type: none"><li>• <b>event_id</b> stored event identifier on VTN</li><li>• <b>N intervals on VTN</b> queryable by all VENs</li><li>• <b>Linked to program_id</b> VENs can filter by program</li><li>• <b>HTTP 201 confirmed</b> raised if creation fails</li></ul>

**CUSTOMIZE**

- Interval duration: "PT1H" (1 hr default) — OpenADR 3 also supports "PT15M" for 15-minute intervals
- Price rounding: round(price, 5) — adjust decimal places to match utility data precision requirements
- Multiple events: publish separate events per tariff zone, VEN group, or building type
- Event expiry: set intervalPeriod.duration to automatically expire stale events on the VTN

26

# Step 5: Read Events as a VEN

## 5 Read Events as a VEN


Simulate the water heater controller reading price signals from the VTN

INPUTS	WHAT IT DOES	OUTPUTS
<ul style="list-style-type: none"><li><b>VEN_CLIENT_ID / SECRET</b> VEN credentials (Step 1)</li><li><b>VTN_BASE_URL</b> same server as Steps 3–4</li><li><b>(Optional) programID</b> filter events by program</li></ul>	<p><b>GET /events (as VEN)</b></p> <pre>Sort intervals by interval id for each payload where payload["type"] == "PRICE":     prices.append(values[0])</pre> <p>Prices now ready for the LP Scheduler (Step 6)</p>	<ul style="list-style-type: none"><li><b>prices [\$/kWh] × N hrs</b> same prices, from VEN perspective</li><li><b>events[] (raw list)</b> full event objects from VTN</li><li><b>event metadata</b> eventName, intervalPeriod.start</li></ul>

**CUSTOMIZE**

- Filter by program: append ?programID=X to GET /events to subscribe to a specific tariff
- Multiple programs: loop events[] and select by evt["programID"] for multi-tariff deployments
- Real deployment: replace this notebook GET with a persistent VEN client process (e.g. openleadr-python)
- Report-back: POST /reports to send HPWH telemetry (SOC, temperatures, energy use) to the VTN

# Step 6: Run LP Scheduler



# Step 7: Generate CTA-2045 Schedule

## 7 Generate CTA-2045 Schedule

Map the HP schedule to discrete demand-response commands for the water heater

INPUTS	WHAT IT DOES	OUTPUTS								
<ul style="list-style-type: none"><li>• <b>schedule dict (Step 6)</b> schedule["control"] per hour</li><li>• <b>params (Step 6)</b> params["max_input"] per hour</li></ul>	<p><b>hpwh_load_shift_to_cta2045(sched, params)</b></p> <p>fraction = control[h] / max_input[h]</p> <table border="1"><tr><td>fraction = 0%</td><td>Shed (-1)</td></tr><tr><td>0% – 30%</td><td>Normal (0)</td></tr><tr><td>30% – 80%</td><td>Load Up (1)</td></tr><tr><td>≥ 80%</td><td>Adv LU (2)</td></tr></table> <p>Repeat for every hour in horizon</p>	fraction = 0%	Shed (-1)	0% – 30%	Normal (0)	30% – 80%	Load Up (1)	≥ 80%	Adv LU (2)	<ul style="list-style-type: none"><li>• <b>signals per hour</b> list of [-1, 0, 1, 2] values</li><li>• <b>signal_names per hour</b> ["Shed", "Normal", ...]</li><li>• <b>Formatted schedule string</b> via format_schedule() for logging</li><li>• <b>plot_schedule() figure</b> bar chart of signals over time</li></ul>
fraction = 0%	Shed (-1)									
0% – 30%	Normal (0)									
30% – 80%	Load Up (1)									
≥ 80%	Adv LU (2)									

**CUSTOMIZE**

- Thresholds: change 30% and 80% in hpwh\_load\_shift\_to\_cta2045() to match your device's response curve
- Price-based mapping: use prices\_to\_cta2045(prices) instead — assigns signals by price percentile
- Physical control: replace format\_schedule() print with an API call to your HPWH hardware
- Fleet deployment: loop over units and call hpwh\_load\_shift() + hpwh\_load\_shift\_to\_cta2045() per device

## **6. How to Use This Software**

# Quick Setup

## 1. Clone and install

```
git clone <repository-url>
cd annex96-a3-hotwater
pip install -r requirements.txt
```

## 2. Start a VTN (choose one)

**Option A: Python VTN** (contact [anandkrp@andrew.cmu.edu](mailto:anandkrp@andrew.cmu.edu) for access)

```
cd openadr3-vtn-reference-implementation
virtualenv venv && source venv/bin/activate
pip install -r requirements.txt && python -m swagger_server
```

**Option B: openleadr-rs** (open source — see `instructions-openleadr.ipynb`)

# Run the Quickstart

## 3. Launch the notebook

```
jupyter notebook quickstart.ipynb          # for Python VTN  
jupyter notebook quickstart-openleadr.ipynb # for openleadr-rs
```

## 4. What the notebook does

Step	Action
1–2	Connect to VTN, fetch dynamic prices from Olivine API
3–4	Create pricing program, publish hourly price event
5	Read events as a VEN
6	Run LP Scheduler and plot optimal schedule
7	Generate CTA-2045 command schedule

# Extending the Software

- **Customize HPWH parameters** — Edit Step 6: tank size, HP capacity, COP, draw profile
- **Use your own price data** — Replace Olivine API with your own source
- **Swap the scheduler** — `hpwh_load_shift_lp` (optimal)  `hpwh_load_shift_heuristic` (no scipy needed)
- **Integrate with CTA-2045 hardware** — Connect generated schedules to physical devices
- **Connect to a real VEN** — Replace notebook HTTP calls with a persistent VEN client

# Resources and References

- **Repository:** annex96-a3-hotwater/
- **OpenADR 3.0.1 Spec:** included in repo, or [openadr.org](http://openadr.org)
- **openleadr-rs:** [github.com/OpenLEADR/openleadr-rs](https://github.com/OpenLEADR/openleadr-rs)
- **Olivine API:** [api.olivineinc.com/i/oe/pricing/signal/paced/etou-dyn](https://api.olivineinc.com/i/oe/pricing/signal/paced/etou-dyn)
- **scipy linprog / HiGHS:** [scipy.org](http://scipy.org) — LP solver used by the scheduler

# Thank You

**Anand Prakash**

[anandkrp@andrew.cmu.edu](mailto:anandkrp@andrew.cmu.edu)

IEA EBC Annex 96 — Activity A3  
Carnegie Mellon University

*All source code and documentation available in the [annex96-a3-hotwater](#) repository*